



Orchestrating the Agentic Defense Pipeline


Moving security AI from single-prompt chat to a deterministic, multi-agent architecture.



[SYSTEM_LOG: 2024-10-27 10:45:32 UTC] | [PIPELINE_VER: v3.5.2-BETA]
[IP_WHITELIST: 18.0.2.15, 172.16.1.8] | [JSON_SCHEMA: DEFINED]



The Attacker: Multi-step, multi-stream, adaptive.




The Defender: A single-completion analyst on a specialty task.


Coordinated attackers compose signals across systems.
The defender working one prompt at a time is structurally outgunned.

Why Single-Prompt Analysis Fails the Defensibility Test


No Cross-Stream Evidence



Hallucinated Correlations



No Defensible Chain




If your tooling can't say "here are the raw log lines that prove it,"
your tooling can't be trusted in the room where decisions get made.

The Architectural Shift


Single-Prompt AI 	Agentic Defense 
Scope of Vision: Sees exactly one artifact at a time.	Scope of Vision: Specialized agents monitor dedicated streams.
Cross-Stream Logic: Blind to bypasses; forces artificial correlations.	Cross-Stream Logic: Composed centrally at exactly one designated layer.
False Positive Risk: High (no native concept of independent signals).	False Positive Risk: Low (typed contracts prevent hallucinated links).
Auditability: Zero. Outputs cannot be walked back to raw logs.	Auditability: Ironclad. Every output carries IDs back to the raw source.

The Composition Layer Earns Its Name




The agent boundary is sharp. Cross-stream correlation lives at exactly one layer—not smeared across every prompt.

Three Streams, One Coordinated Attack



The Load-Bearing Negative Check



The two API token-expiration 401s **MUST NOT** be correlated with the password spray.
Refusing to invent unrelated correlations is the quiet win that earns trust.

Orchestration Complete

This is a password spray from 185.220.101.42
targeting 30 users.

Block the IP at the edge.


Audit the affected user accounts.

Ignore the unrelated API token expirations.

┌ The spine of the system. ┐
└ The mechanism to explicitly ┘
dismiss noise. ┘




The Defensibility Chain



Every artifact carries IDs back to its source. Click any action; walk the chain back to the literal raw log line. No extra API calls. No model recollection. Just a graph walk.

Every Investigation, Perfectly Replayable



Audit

Six months later, a finding is challenged. The full evidence chain is reconstructible.


Reproducibility

Same input, same provenance graph, every time. Bugs are findable.

Trust at Scale


When producing dozens of findings a week, full provenance makes spot-checking viable.

Trust Earned, Not Assumed



The human elevates from manual orchestrator to auditor. The AI executes the the loop; the human decides what to do about the findings.

Run the Architecture in Your Browser



Phase 1 (Manual Walkthrough)

Copy stream prompts, paste into AI, paste JSON back. Expectation panel evaluates in plain English in real time.

Phase 2 (Automated Execution)

Drop in an API key. Click run. Model returns a JSON finding in ten seconds across all streams.

One self-contained HTML file. In 15 minutes, see exactly why a multi-agent pipeline catches what single-prompt scanners cannot.